

Q 2.4.17. Standard Gradient Descent vs. Natural Gradient Descent

In this notebook, we compare standard gradient descent (GD) and natural gradient descent (NGD) for estimating the parameters of a 1D Gamma distribution from observed data. We will generate synthetic data from a known Gamma distribution and then attempt to recover its shape and scale parameters using both optimization methods.

▼ 1. Configuration

We define the true parameters of the Gamma distribution, the number of data points to generate, the learning rate, the number of epochs for optimization, and the initial guesses for the parameters. We also set a random seed for reproducibility.

```
import torch
import numpy as np
import matplotlib.pyplot as plt
from torch.distributions import Gamma
# -----

# True parameters of the Gamma distribution we want to discover
ALPHA_TRUE = 3.0    # shape
BETA_TRUE   = 2.0    # scale

# Number of observed data points
N_DATA = 1000

# Optimization parameters
LEARNING_RATE = 0.1
EPOCHS = 150

# Initial "wrong" guess for our parameters (still positive)
ALPHA_INIT = 0.5
BETA_INIT   = 8.0

# Fix random seed for reproducibility (optional)
torch.manual_seed(0)
```

```
<torch._C.Generator at 0x122105750>
```

▼ 2. Data Generation

We generate synthetic data from the true Gamma distribution using PyTorch's `Gamma` distribution class. We sample `N_DATA` points and print the sample mean and variance to verify the generated data.

```
# Our parameterisation is shape-scale, but PyTorch's Gamma uses sha
rate_true = 1.0 / BETA_TRUE
dist_true = Gamma(concentration=torch.tensor(ALPHA_TRUE),
                  rate=torch.tensor(rate_true))

data = dist_true.sample((N_DATA,))

print(f"Generated {N_DATA} data points from Gamma(alpha={ALPHA_TRUE})
print(f"Sample mean:      {data.mean().item():.4f}")
print(f"Sample variance: {data.var().item():.4f}\n")
```

```
Generated 1000 data points from Gamma(alpha=3.0, beta=2.0)
Sample mean:      6.0692
Sample variance: 11.9481
```

▼ 3. Loss function & parameter init

We define the negative log-likelihood loss function for the Gamma distribution. We also initialize the parameters for both standard gradient descent and natural gradient descent, and set up history trackers to record the parameter values over epochs.

```
# ToDo: Loss function
def gamma_nll(alpha, beta, data_points):
    """
    ToDo:
        Implement the average negative log-likelihood for Gamma dis

    Hints:
        - Enforce positivity using clamp (e.g. min=1e-4).
        - PyTorch's Gamma takes (concentration=alpha, rate=1/beta).
        - Return the *mean* negative log-likelihood.

    .....
    alpha = torch.clamp(alpha, min=1e-4)
    beta = torch.clamp(beta, min=1e-4)

    # Convert data_points to tensor if not already
    data_tensor = torch.tensor(data_points) if not torch.is_tensor(
        data_points) else data_points

    empirical_mean_x = data_tensor.mean()
    empirical_mean_log_x = torch.log(data_tensor).mean()
    nll = +alpha*torch.log(beta) + torch.lgamma(alpha) - (alpha-1)*
    return nll
```

```
# Parameters for Standard Gradient Descent (GD)
alpha_gd = torch.tensor(ALPHA_INIT, requires_grad=True)
beta_gd = torch.tensor(BETA_INIT, requires_grad=True)

# Parameters for Natural Gradient Descent (NGD)
alpha_ngd = torch.tensor(ALPHA_INIT, requires_grad=True)
beta_ngd = torch.tensor(BETA_INIT, requires_grad=True)

# History trackers
history_gd = []
history_ngd = []
history_loss_gd = []
history_loss_ngd = []
```

▼ 4. Fisher Information inverse

```

def fisher_inverse(alpha, beta):
    """
    TODO:
        Implement the inverse Fisher Information matrix  $F^{-1}(\alpha, \beta)$ 
        for the Gamma(shape= $\alpha$ , scale= $\beta$ ) distribution.

    Theory:
         $F(\alpha, \beta) =$ 
        
$$\begin{bmatrix} \psi_1(\alpha) & 1/\beta \\ 1/\beta & \alpha/\beta^2 \end{bmatrix}$$


         $F^{-1}(\alpha, \beta) =$ 
        
$$\frac{1}{\alpha \psi_1(\alpha) - 1} \begin{bmatrix} \alpha & -\beta \\ -\beta & \beta^2 \psi_1(\alpha) \end{bmatrix}$$


    Hints:
        - Use torch.polygamma(1, alpha) for  $\psi_1(\alpha)$  (trigamma).
        - Make sure to detach alpha, beta so  $F^{-1}$  is not part of the graph.
    """
    alpha_detached = alpha.detach()
    beta_detached = beta.detach()

    external_factor = 1/(alpha_detached * torch.polygamma(1, alpha))
    inv11 = alpha_detached * external_factor
    inv12 = -beta_detached * external_factor
    inv22 = beta_detached**2 * torch.polygamma(1, alpha_detached)
    return inv11, inv12, inv22

```

4. Optimization Loop

We run the optimization loop for a specified number of epochs. In each epoch, we perform both standard gradient descent and natural gradient descent updates. We compute the gradients, build the Fisher Information Matrix for NGD, and update the parameters accordingly. We also log the parameter values and losses at regular intervals.

Inizia a programmare o genera codice con l'IA.

```

print(f"Optimizing with LR={LEARNING_RATE} for {EPOCHS} epochs...")

for epoch in range(EPOCHS):

```

```
# ===== A. Standard Gradient Descent (GD) =====

if alpha_gd.grad is not None:
    alpha_gd.grad.zero_()
if beta_gd.grad is not None:
    beta_gd.grad.zero_()

loss_gd = gamma_nll(alpha_gd, beta_gd, data)

loss_gd.backward()

with torch.no_grad():
    alpha_gd -= LEARNING_RATE * alpha_gd.grad
    beta_gd -= LEARNING_RATE * beta_gd.grad

    alpha_gd.clamp_(min=1e-4)
    beta_gd.clamp_(min=1e-4)

history_loss_gd.append(loss_gd.item())
history_gd.append((alpha_gd.item(), beta_gd.item()))

# ===== B. Natural Gradient Descent (NGD) =====

if alpha_ngd.grad is not None:
    alpha_ngd.grad.zero_()
if beta_ngd.grad is not None:
    beta_ngd.grad.zero_()

loss_ngd = gamma_nll(alpha_ngd, beta_ngd, data)
loss_ngd.backward()

g_alpha = alpha_ngd.grad
g_beta = beta_ngd.grad

# ToDo : compute natural gradient using  $F^{-1}(\alpha, \beta)$ 
# 1) Get  $F^{-1}$  entries using fisher_inverse(...)
# 2) Compute:
#     - ng_alpha
#     - ng_beta

inv11, inv12, inv22 = fisher_inverse(alpha_ngd, beta_ngd)

ng_alpha = inv11*g_alpha + inv12*g_beta
ng_beta = inv12*g_alpha + inv22*g_beta

with torch.no_grad():
    alpha_ngd -= LEARNING_RATE * ng_alpha
```

```
        beta_ngd -= LEARNING_RATE * ng_beta

        alpha_ngd.clamp_(min=1e-4)
        beta_ngd.clamp_(min=1e-4)

    history_loss_ngd.append(loss_ngd.item())
    history_ngd.append((alpha_ngd.item(), beta_ngd.item()))

    if (epoch + 1) % 15 == 0 or epoch == 0:
        print(f"\n--- Epoch {epoch + 1} ---")
        print(f"  GD: alpha={alpha_gd.item():.4f}, beta={beta_gd.i
              f"Loss={loss_gd.item():.4f}"")
        print(f"  NGD: alpha={alpha_ngd.item():.4f}, beta={beta_ngd
              f"Loss={loss_ngd.item():.4f}"")\n\nOptimization finished."
```

Optimizing with LR=0.1 for 150 epochs...

--- Epoch 1 ---

GD: alpha=0.6516, beta=8.0032, Loss=3.1866
NGD: alpha=0.5340, beta=7.8695, Loss=3.1866

--- Epoch 15 ---

GD: alpha=1.0504, beta=7.9743, Loss=2.8331
NGD: alpha=1.2770, beta=4.3331, Loss=2.7370

--- Epoch 30 ---

GD: alpha=1.0853, beta=7.9160, Loss=2.8296
NGD: alpha=2.3102, beta=2.5556, Loss=2.5688

--- Epoch 45 ---

GD: alpha=1.0933, beta=7.8551, Loss=2.8270
NGD: alpha=2.8686, beta=2.1013, Loss=2.5450

--- Epoch 60 ---

GD: alpha=1.0989, beta=7.7937, Loss=2.8245
NGD: alpha=3.0273, beta=2.0018, Loss=2.5435

--- Epoch 75 ---

GD: alpha=1.1045, beta=7.7319, Loss=2.8219
NGD: alpha=3.0627, beta=1.9810, Loss=2.5434

--- Epoch 90 ---

GD: alpha=1.1101, beta=7.6698, Loss=2.8193
NGD: alpha=3.0701, beta=1.9768, Loss=2.5434

--- Epoch 105 ---

```

GD: alpha=1.1159, beta=7.6072, Loss=2.8167
NGD: alpha=3.0716, beta=1.9759, Loss=2.5434

--- Epoch 120 ---
GD: alpha=1.1217, beta=7.5444, Loss=2.8141
NGD: alpha=3.0719, beta=1.9757, Loss=2.5434

--- Epoch 135 ---
GD: alpha=1.1277, beta=7.4811, Loss=2.8114
NGD: alpha=3.0720, beta=1.9757, Loss=2.5434

--- Epoch 150 ---
GD: alpha=1.1339, beta=7.4175, Loss=2.8086
NGD: alpha=3.0720, beta=1.9756, Loss=2.5434

Optimization finished.

```

▼ Q 2.4.18. Plotting Results

To illustrate the difference between standard gradients and natural gradients, we print out the gradients computed in the first epoch for both methods.

```

hist_gd_np = np.array(history_gd)
hist_ngd_np = np.array(history_ngd)
hist_loss_gd = np.array(history_loss_gd)
hist_loss_ngd = np.array(history_loss_ngd)

fig, (ax1, ax2, ax3) = plt.subplots(3, 1, figsize=(12, 10), sharex=True)
fig.suptitle(f"Gamma: Standard Gradient vs. Natural Gradient "
             f"(LR={LEARNING_RATE}, N={N_DATA})", fontsize=16)

# Plot 1: alpha (shape)
ax1.plot(hist_gd_np[:, 0], label="GD alpha", color='blue', linestyle='solid')
ax1.plot(hist_ngd_np[:, 0], label="NGD alpha", color='red', linestyle='solid')
ax1.axhline(ALPHA_TRUE, color='black', linestyle='dashed', label=f"True value")
ax1.set_ylabel("Shape parameter $\alpha$")
ax1.legend()
ax1.grid(True)

# Plot 2: beta (scale)
ax2.plot(hist_gd_np[:, 1], label="GD beta", color='blue', linestyle='solid')
ax2.plot(hist_ngd_np[:, 1], label="NGD beta", color='red', linestyle='solid')
ax2.axhline(BETA_TRUE, color='black', linestyle='dashed', label=f"True value")
ax2.set_xlabel("Epoch")

```

```

ax2.set_ylabel("Scale parameter $\beta$")
ax2.legend()
ax2.grid(True)

# Plot 3: Negative Log-Likelihood (NLL) Loss
ax3.plot(hist_loss_gd, label="GD NLL Loss", color='blue', linestyle='solid')
ax3.plot(hist_loss_ngd, label="NGD NLL Loss", color='red', linestyle='solid')
ax3.set_xlabel("Epoch")
ax3.set_ylabel("NLL Loss")
ax3.legend()
ax3.grid(True)

plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()

```

Gamma: Standard Gradient vs. Natural Gradient (LR=0.1, N=1000)



